LIBERO-LoHo | Astra vs Opus

5 saved runs per model/task | mean ±1 sample SD

GPT-6 Astra max Opus 5.5 max

Goal completion score

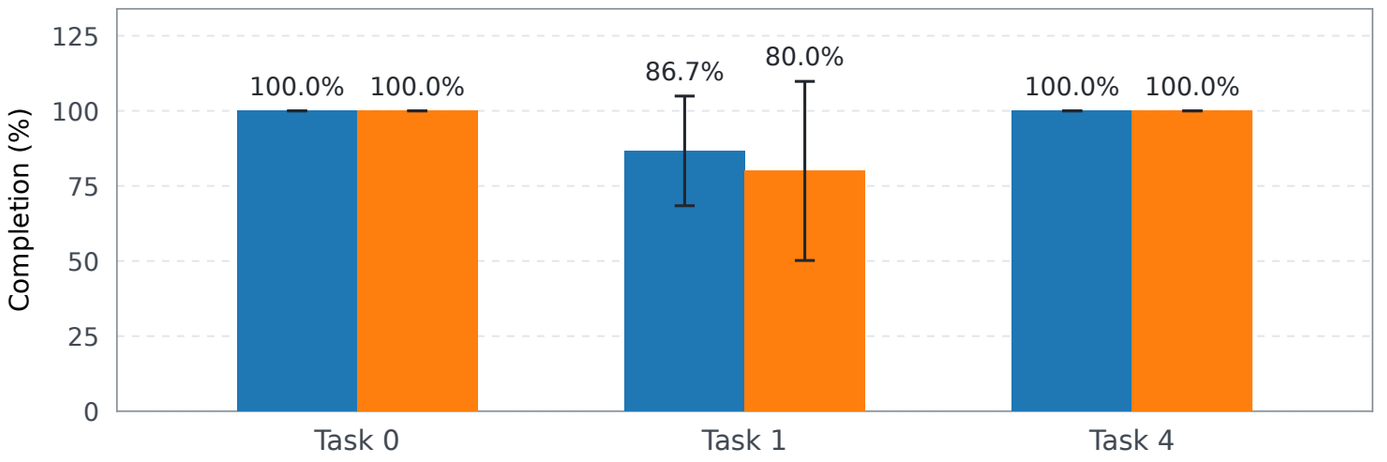

Total tokens and estimated API cost

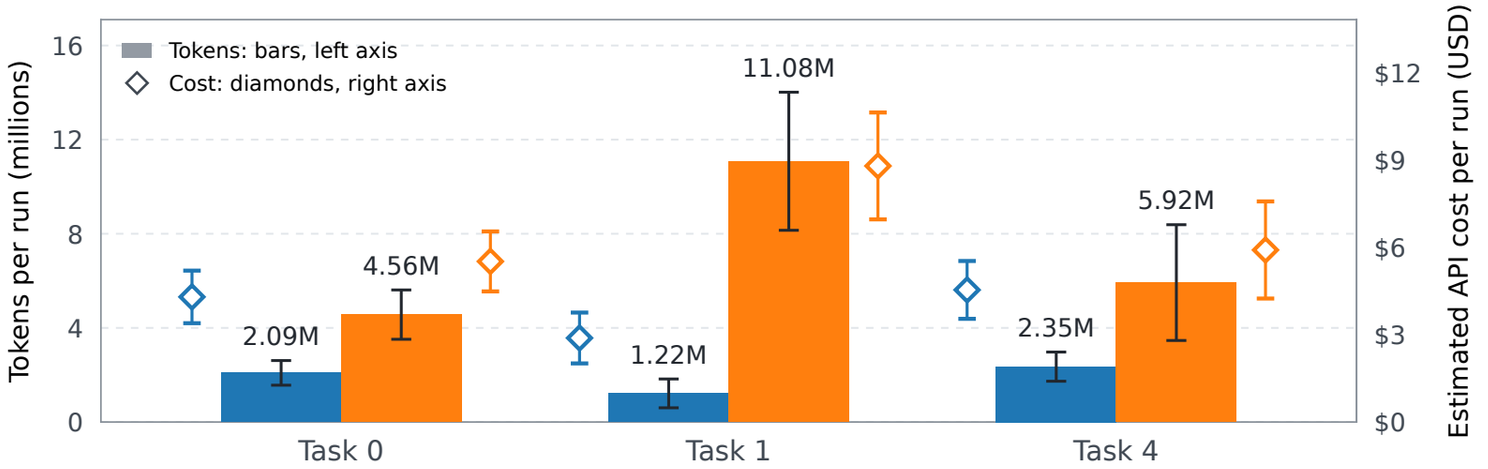

Wall-clock time

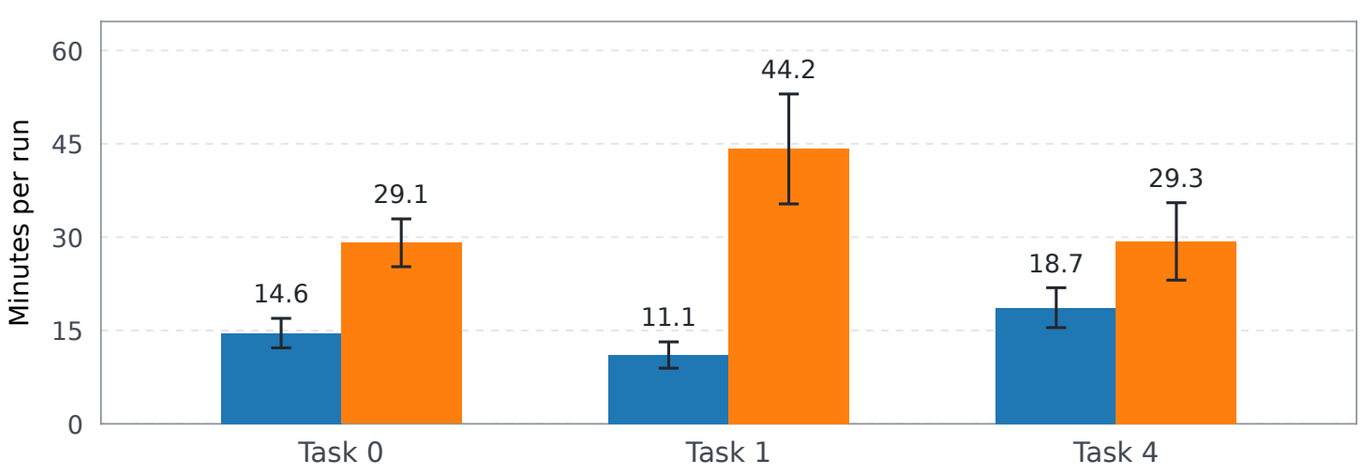

Tokens: input + output, including cache usage and reported reasoning.
Whiskers: sample SD, not confidence intervals. Task 4: retained run set.
All runs: state 0, seed 0, max effort, static + wrist RGB and proprioception.