7 | Compact results table

Mean ± sample standard deviation across five saved episodes per model/task

Task	Model	Full successes	Goal completion (%)	Output tokens (k/run)	Estimated API cost (USD/run)	Wall-clock time (min/run)
0	Astra	5/5	100.0 ± 0.0	27.9 ± 5.7	4.31 ± 0.90	14.6 ± 2.4
0	Opus	5/5	100.0 ± 0.0	174.9 ± 28.1	5.53 ± 1.03	29.1 ± 3.8
1	Astra	3/5	86.7 ± 18.3	21.3 ± 3.0	2.89 ± 0.87	11.1 ± 2.1
1	Opus	3/5	80.0 ± 29.8	257.9 ± 52.3	8.81 ± 1.84	44.2 ± 8.8
4	Astra	5/5	100.0 ± 0.0	26.5 ± 5.9	4.55 ± 0.99	18.7 ± 3.2
4	Opus	5/5	100.0 ± 0.0	178.3 ± 36.2	5.92 ± 1.67	29.3 ± 6.2

Output tokens include reasoning; costs are public-rate estimates; wall-clock includes API waiting.
Full successes are counts, not mean ± SD. Task 4 uses the current retained run set; all episodes use state 0 and seed 0.